


# Intro to Kubernetes

## for developers



April 8, 2024

ZHealth Documentation

# Background

+

- Docker: A Project-Based Approach to Learning by Jason Cannon
- The Kubernetes Book: 2023 Edition by Nigel Poulton
- LFS158x: Intro to Kubernetes
- Documentation at [kubernetes.io](https://kubernetes.io)



# What is Kubernetes?

+

- A way to automate the management of distributed services at scale running on unreliable hardware.
- Manual administration requires understanding the current state in order to make changes.
- This breaks down at scale: nobody can keep track of all the machines, services, settings, dependencies, versions, and cloud infrastructure as the number of services increases in orders of magnitude.

## What does it do?

+

- You tell the orchestrator your desired end state.
- The orchestrator applies changes based on the current state that it's able to keep track of because it's an automated and distributed service itself.
- If something goes wrong, it's easier to isolate the problem down to a transactional interaction of a few components, so that it can be reproduced and fixed.
- Without this structure, problems are band-aided and Production team is endlessly fighting fires.

## How is it used?

+

- Kubernetes builds on the concept of a container.
- It lets you deploy a cluster of containerized services by using a manifest.
- The manifest is a declaration of which containers you want to be running along with their dependencies.
- Kubernetes has a package manager called Helm that can be used to assemble a manifest from standard parts.


## How does it work?



- Kubernetes reads the given manifest and continuously applies changes to the underlying infrastructure, services, and configuration.
- Kubernetes talks to the underlying infrastructure via plugins, such as the container runtime, storage, software-defined networking, and cloud services like load balancers and firewalls.
- Manifest changes are applied via dashboard or CLI.
- Manifests can be source controlled & applied by CI.

# What's in a cluster?



+



# A cluster of clusters?

+

- Kubernetes Federation lets you manage clusters in multiple availability zones to enable failover.





# Kubernetes History



- Kubernetes was created by Google to support Search and Ads products responsible for most of its revenue.
- Both products required running a large number of jobs on cheap hardware with a free operating system, all Google could afford when they were starting out.
- Kubernetes was released to steer developers toward Google and poach customers from Amazon and Microsoft.
- Amazon and Microsoft promptly released plugins that only work with Kubernetes on their platforms.

# Why use Kubernetes?



- Standardizes cloud provider features.
- Simplifies deploying a new version of a service.
- Unhealthy services are automatically restarted.
- Unhealthy machines are taken offline and services are re-scheduled on different machines automatically.
- Simplifies dynamically assembling infrastructure and services taking all dependencies into account.
- Canary and Blue/Green deployments are supported out of the box.

## Cons

+

- The complexity of maintaining large distributed systems have simply been shifted to another frequency.
- Instead of struggling with network administration, now you struggle with writing manifests and configuring plugins so that orchestrator can apply changes.
- Complicated enough that you can't muddle your way through it for real use in production.
- Another package ecosystem to learn.

# Kubernetes in Practice

+

- Security is complicated to setup properly, with lots of conflicting constraints to keep in mind.
- High-availability/multi-zone deployment requires Kubernetes Federation.
- Software-defined networking was built on top of Linux, understanding limitations requires sysadmin experience.
- You have to memorize what each feature can handle.
- Plugin system for cloud-provider-specific services ensures learning "never stops".

## Growing pains



- Some features like load balancing and software-defined networking were meant to be a core part of the system, but after the public release it became apparent that they didn't handle many real-world scenarios.
- In the end Kubernetes became "a little of everything for everybody", with muddy "design by committee".
- Some features superseded by plugins but left in for backward compatibility.

## Should we use it?

+

- Maintaining services at a scale: we are not yet affected by issues of scale.
- Reducing cloud hosting costs: we won't realize cost savings unless we switch to a different cloud.
- Running on our own hardware: running on managed cloud is cheaper without a production team.
- Our deployments are manageable, can't afford canary & blue/green yet due to cloud infrastructure costs.